System One models — complete comparison table

The same model-by-model inventory as the preceding answer, formatted for PDF with clickable source links.

Columns included: model, availability (API/weights), text/image/audio/video support, source implementation, available formats, and on-device deployment potential.

Evidence labels: Available — source explicitly identifies the artifact; Not verified — exact artifact not established, which does not prove impossibility; Not documented — no reliable confirmation of the modality. “Possible” does not mean NPU deployment is proven.

A. Hosted API models

These are the named commercial or hosted decision models listed in the current model directory. Some also publish weights, so API does not always mean API only.

Model	API available?	Downloadable weights?	Text	Image	Audio	Video	Source implementation available?	Formats verified	On-device potential
Jev — TypeSafe AI	Yes	Not identified	Yes	Not verified	Not verified	Not verified	API/SDK; model internals not public	Hosted API only identified	No direct path identified
Tev1 — Together AI	Yes, early access	Experimental/open variant status needs checking	Yes	Not verified	Not verified	Not verified	API integration	Hosted API; experimental variant has separate artifacts	Only if exact open variant is obtained
Solar Decide — Upstage	Yes, early access	Not identified	Yes	Not verified	Not verified	Not verified	API/SDK	Hosted API	No direct path identified
Span-01 — Respan	Yes	Not identified	Yes; agent-trace decisions	Not verified	Not verified	Not verified	API/SDK	Hosted API	No direct path identified
d1 — Liquid AI	Yes	Yes; separate d1-3B and d1-omni-600M variants	Yes	Yes in open variants	Audio in d1-omni-600M	Not verified for standard d1	API and open-variant implementation	Hosted API; separate open-weight checkpoints	Possible using open variants
d1-ultra — Liquid AI	Yes	Not identified	Yes	Yes reported	Not verified	Not verified	API/SDK	Hosted API	No direct path identified
GLiDE — Fastino Labs	Yes	No public weights verified for this hosted model	Yes	Not verified	Not verified	Not verified	API/SDK	Hosted API	No direct path identified
Decider 1 — meraGPT	Yes	Separate open Decider models exist; same weights not established	Yes	Not verified for API model	Not verified	Not verified	API/SDK; separate open implementation exists	Hosted API; open Decider variants have separate formats	Use an open variant; do not assume API weights are downloadable
Mercury Decide — Inception	Yes	Not identified	Yes	Not verified	Not verified	Not verified	API/SDK	Hosted API	No direct path identified
Clef — Cloudflare	Yes	Yes	Yes	Yes	Not verified for standard Clef	Not verified for standard Clef	Public model materials and API	Open weights; exact export formats need variant-level verification	Possible; local export and hardware testing needed
Clef-flash — Cloudflare	Yes	Yes	Yes	Yes	Not verified	Not verified	Public model materials and API	Open weights; no universal NPU artifact established	Possible; local export and hardware testing needed
Clef-omni — Cloudflare	Yes	Yes	Yes	Yes	Yes	Yes	Public model materials and API	Open weights; exact portable export not verified	Possible, but multimodal deployment is more complex
pplx-decider — Perplexity	Yes	Yes	Yes	Yes	Not verified	Not verified	Public weights and API integration	Open weights; exact ONNX/Core ML export not verified	Possible; local export and hardware testing needed

Model	API available?	Downloadable weights?	Text	Image	Audio	Video	Source implementation available?	Formats verified	On-device potential
Kai — Hanzo AI	Yes	Not identified	Yes	Not verified	Not verified	Not verified	API/SDK	Hosted API	No direct path identified
Drex — Nace.AI	Yes	Yes; separate Drex 1.5 weights reported	Yes	Not verified	Not verified	Not verified	Public implementation for open variant; hosted API	Python/vLLM and GGUF/llama.cpp paths reported for variants	Possible with open variant; validate license and runtime
Celeris-1 Decision — Celeris	Yes	Not identified	Yes	Yes	Not verified	Not verified	API/SDK	Hosted API	No direct path identified
Microsoft-Decision-1	Yes	No	Yes	No; documented as text-only	No verified support	No verified support	API/SDK; model weights closed	Hosted API	No direct path identified
OpenAI Decisions API	Yes; public beta reported in directory	No downloadable weights identified	Yes	Image input reported	Not verified	Not verified	API/SDK; model internals not public	Hosted API	No direct path identified

Links: Jev documentation · System One API · Jev alternatives guide · Model directory · OpenAI API docs

B. Open/hosted registry models

The registry includes base models, fine-tunes and hardware-oriented variants. An export is only listed where the exact artifact was established; otherwise it is marked not verified.

Model	API?	Weights?	Text	Image	Audio	Video	Source implementation	Formats verified	On-device potential
Vega	Not confirmed	Yes	Yes	Not verified for exact entry	Not verified	Not verified	Model card; full source status needs checking	Exact export not verified	Possible; verify checkpoint and vision export
AJEV	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Model card; source completeness needs checking	Exact export not verified	Possible, but 26B is a large edge target
Nativ-HE-Decision	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Public listing; code completeness needs checking	Exact export not verified	Possible; compact 364M; validate Hebrew fit
Strom	Yes	No public weights identified	Yes	Not verified	Not verified	Not verified	API; model internals not verified	Hosted API	No direct path identified
Canopy-Jev	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Model card; source completeness needs checking	Exact export not verified	Possible, but 27B is large for edge
Docto-Decision	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Public listing; code completeness needs checking	Exact export not verified	Possible; 4.7B needs memory/latency tests
Gutsy	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Listing; source completeness needs checking	Q8_0 run reported; other exports not verified	Promising size; test runtime and accuracy
ModernJev-Decision	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Listing; source completeness needs checking	Exact export not verified	Promising compact 150M candidate
StartLux-Decision	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Listing; source completeness needs checking	Exact export not verified	Possible; 4.7B needs memory/runtime tests
Bekko-System-One	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Listing; code completeness needs checking	Exact export not verified	Promising size; verify artifacts
BioDecision	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Public listing; source completeness needs checking	Exact export not verified	Possible; domain-specific 4B model
Bev-Decider	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Listing; code completeness needs checking	Exact export not verified	Possible; compact 478M candidate
Rizzo-Flow	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Listing; code completeness needs checking	Q8_0 run reported; other exports not verified	Possible; 4B; validate local runtime
Security-One	Yes	Yes	Yes	Not verified	Not verified	Not verified	Public listing and hosted API	Exact export not verified	Possible, but 27B is large for edge
GLiNER2.5-Decide	Yes	Yes	Yes	Not verified for checkpoint	Not verified	Not verified	Public weights and implementation reported	Exact ONNX export not verified here	Possible; verify NPU operators
Decider	Hosted variants exist	Yes	Yes	Separate vision build reported	Not verified	Not verified	Public implementations and weights	vLLM/Python for some variants; ONNX varies	Possible; variant-specific
Julia-1	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Model card/local implementation	GGUF mirror reported; exact ONNX not verified	Promising compact 144M candidate
CLM	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Public implementation and weights reported	PyTorch/source path; exact ONNX not verified	Possible; 8B needs memory/quantization tests
NeoHorse-Jev	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Listing; source completeness needs checking	vLLM run reported; portable export not verified	Possible; 4B; validate target runtime
Lumma-FEV	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Listing; source completeness needs checking	Exact export not verified	Possible; 649M; verify formats

Model	API?	Weights?	Text	Image	Audio	Video	Source implementation	Formats verified	On-device potential
Kev	Hosted options exist	Yes	Yes	Not verified for standard checkpoint	Not verified	Not verified	Public implementation/training code reported	CUDA, ROCm and MLX paths reported; ONNX varies by variant	Strong local candidate; NPU still unverified
Laya	Not required for local use	Yes	Yes	Not verified for standard checkpoint	Not verified	Not verified	Public Python implementation	ONNX, Core ML and MLX community paths; GGUF depends on port	Strong first export candidate; validate target NPU
Open-Jev-DeBERTa-v3-large	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Public listing and implementation	Exact ONNX export not verified for checkpoint	Possible; 434M encoder candidate
Mira	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Listing; full source status needs checking	Exact export not verified	Possible; confirm size and formats
Caliber-Flex-1B	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Listing; code completeness needs checking	Exact export not verified	Possible; verify size, license and export
Caliber-Flex-5B	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Listing; code completeness needs checking	Exact export not verified	Possible; higher memory needs
Laya Typed-Decisions NoulXP	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Variant/package listing	NoulXP package; portable export needs checking	Possible; test package runtime/backend
Laya NoulXP-GPU	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Variant/package listing	GPU-oriented package	Possible on supported GPU; no automatic NPU support
Laya NoulXP-CPU	Not confirmed	Yes	Yes	Not verified	Not verified	Not verified	Variant/package listing	CPU-oriented package	CPU baseline; NPU acceleration not established
Laya Multilingual NoulXP	Not confirmed	Yes	Yes; multilingual	Not verified	Not verified	Not verified	Variant/package listing	NoulXP package; portable export needs checking	Possible; test language accuracy and backend

Registry: System One model registry · Laya listing · Kev listing

C. Additional open implementations and multimodal variants

These are separate implementations or variants identified in the open-implementation directory. They are listed individually because they may offer practical local deployment routes.

Model / implementation	API?	Weights?	Text	Image	Audio	Video	Source implementation	Formats verified	On-device potential
AgentJev-0.6B	Self-hosting/API serving reported	Yes	Yes	Not verified	Not verified	Not verified	Yes	Python; exact ONNX not verified	Possible; small local candidate
Dohnuts 0.8B	Self-hosting reported	Yes	Yes	Yes	Not verified	Not verified	Yes; training/inference/evaluation toolkit	Weights/Python; exact ONNX not verified	Possible; test multimodal pipeline
Drex 1.5	Yes	Yes	Yes	Not verified	Not verified	Not verified	Yes	Python/vLLM and GGUF/llama.cpp paths reported for variants	Possible; validate license and NPU
FRIDA-Decisions	Self-hosting reported	Yes	Yes	Not verified	Not verified	Not verified	Yes	Torch, int8 ONNX, MLX and vLLM reported	Strong format flexibility; NPU needs operator validation
GUI-Decisions	Self-hosting reported	Yes	Textplus GUI/action representations	GUI grounding; raw-image interface needs checkpoint verification	Not verified	Not verified	Yes; LoRA training code/weights	Model-specific; exact ONNX not verified	Possible; 2B/31B variants target different hardware
JevEmbed	Self-hosting reported	Yes; collection reported	Yes	Not verified	Not verified	Not verified	Yes	Hugging Face collection; exact export not verified	Possible; depends on embedder and decision head
JevK5	Yes; local API reported	Yes	Yes	Not verified	Not verified	Not verified	Yes	Exact ONNX/Core ML not verified	Possible; 4B/9B variants need testing
Jevlike	Self-hosting reported	Yes / code available	Yes	Not verified	Not verified	Not verified	Yes	Python implementation; graph export not verified	Possible; verify configuration and hardware
JevOS	Yes; local server	Yes	Yes; No ul-only serving described	Not verified	Not verified	Not verified	Yes	GGUF and llama.cpp CPU path reported	CPU route reported; NPU not established
Jiwo	Self-hosting reported	Yes	Yes	Not verified	Not verified	Not verified	Yes	Weights; exact ONNX not verified	Possible; 0.8B/4B variants need testing
NanoJev	Self-hosting reported	Yes	Yes	Not verified	Not verified	Not verified	Yes; dataset/training pipeline reported	Exact ONNX/Core ML not verified	Promising compact candidate; verify runtime
Bespoke Nimble	API also reported	Yes	Yes	Not verified	Not verified	Not verified	Yes; model/training recipe reported	Apple MLX and NVIDIA execution reported; ONNX not verified	Possible; NPU unverified
Open Jev	Self-hosting reported	Yes	Yes	Not verified for compact encoder	Not verified	Not verified	Yes	Python/weights; exact ONNX not verified	Promising 150M encoder candidate
Valen	Self-hosting reported	Yes	Yes	Yes	Not verified	Yes	Yes; code, data pipeline and training/evaluation materials	Exact ONNX not verified	Possible; multimodal export more involved
vegaml	Self-hosting reported	Yes	Yes	Yes	Not verified	Not verified	Yes; code/training/evaluation materials reported	Exact ONNX not verified	Possible; 800M/4B options; validate vision pipeline
WaterSheep	Yes; local API reported	Yes	Yes	Not verified	Not verified	Not verified	Yes; code and weights reported	ONNX build reported for browser inference	Browser-ONNX path reported; native NPU unverified

Directory: System One open implementations · Curated model/runtime census

D. What the format evidence actually tells us

Format or implementation	Models with evidence in reviewed sources	What that establishes	What it does not establish
ONNX	Laya community ports; FRIDA-Decisions; WaterSheep browser build	A graph-format path exists for those particular projects	That every model can export, or that the graph runs on your NPU
Core ML	Laya Core ML port	A local Apple-runtime path exists	Compatibility with Android or another vendor's NPU
MLX	Laya and Kev; FRIDA-Decisions	Local Apple Silicon execution path	Generic NPU portability
GGUF / llama.cpp	JevOS; some Drex variants; other project-specific mirrors	Local CPU-oriented or compatible-backend path	Native acceleration on an arbitrary NPU
PyTorch / Python	Multiple open implementations, including FRIDA-Decisions and model-specific repositories	A reference or local implementation path exists	That export will be lossless or that the NPU supports all operators
Vendor-specific compiled artifact	Laya/Axera community port listed	A model can be adapted to a particular accelerator toolchain	That the same build works on a different NPU

Sources: Open implementations · Curated model/runtime census · DecisionEval model directory

E. Shortlist for your NPU deployment research

Priority	Candidate	Why investigate it	Main unresolved issue
1	Laya	Open weights and multiple local conversion paths are reported	Whether the full graph and preprocessing work on your exact NPU
2	FRIDA-Decisions	Explicit Torch, int8 ONNX, MLX and vLLM paths	Whether the target NPU supports graph operators and shapes
3	WaterSheep	A browser ONNX build is reported	Whether it can be adapted to a native NPU backend
4	Kev	Open weights, source implementation and multiple local execution paths	Whether a suitable export exists for your target accelerator
5	Julia-1	Small 144M model; GGUF mirror reported	Exact runtime compatibility and decision quality
6	Open-Jev-DeBERTa-v3-large	Compact encoder architecture	Export path and compatibility need direct confirmation
7	Clef / Clef-flash	Open weights and image-capable variants	More complex graph; exact export and NPU compatibility
8	d1-omni-600M	Open multimodal variant with image/audio support reported	Exact local format, modality pipeline and NPU compatibility
9	Valen	Public multimodal implementation and training pipeline	More work to deploy image/video path locally
10	NanoJev	Compact open implementation with a reported training pipeline	No verified target-specific NPU export identified

Final interpretation

- If you need text-only routing across 20–30 options, start with Laya, FRIDA-Decisions, Julia-1 and Kev.
- If you need images, add Clef, d1-3B, Valen or vegaml, but verify the exact checkpoint and image preprocessing.
- If you need audio or video, the clearest named candidates in the reviewed sources are d1-omni-600M for image/audio and Clef-omni for image/audio/video. Their exact on-device export remains a separate question.
- If a model is hosted-only and has no downloadable weights, treat exact-model on-device deployment as blocked unless the vendor releases weights or an authorized local implementation.

For your project, keep two separate statuses in the engineering tracker: “local inference demonstrated” and “target NPU acceleration demonstrated.” A model should not be marked NPU-deployable until the second status is proven on the actual device.

Main source links

- Jev documentation
- System One API reference
- System One model registry

- Lava listing
- Key listing
- Open implementations directory
- Curated model/runtime census